

Simulating Time-Varying Performance

Andrew Glover

2023-10-11

Libraries

```
library(ggplot2) # for graphing
library(patchwork) # to add graphs together
library(tibble) # tibbles
```

Functions used for analysis

```

params_mat <- function(k_1, tau_1, k_2, tau_2, change_days=NULL, days) {
  if (length(k_1) == length(tau_1) &&
      length(k_1) == length(tau_2) &&
      length(k_1) == length(k_2) &&
      length(k_1) == length(change_days)+1
  ) {}
  else {stop("check length of parameters")}
  out_matrix <- matrix(0, nrow = days, ncol = 4)
  colnames(out_matrix) <- c("k_1", "tau_1", "k_2", "tau_2")
  bound_1 <- 1; bound_2 <- days
  j <- 0 # counter for index of k_1, tau_1, etc
  for (elem in c(change_days, days)) {
    j <- j + 1
    bound_2 <- elem
    for (i in bound_1:bound_2) {
      out_matrix[i, ] <- c(k_1[[j]], tau_1[[j]], k_2[[j]], tau_2[[j]])
    }
    bound_1 <- elem
  }
  return(out_matrix)
}

```

```
#' Title
#'  
#'  
#'  
#'  
#'  
#'  
#'  
#'  
#'  
#'
```

```

#'
#' @examples
perf_tv <- function(p_0, params_mat, training_load) {
  days <- nrow(params_mat)
  perf_out <- c(rep(NA, days))
  T_1 <- 0; T_2 <- 0
  for (i in 1:days) {
    T_1 <- exp(-1/params_mat[i, "tau_1"])*T_1 + training_load[[i]]
    T_2 <- exp(-1/params_mat[i, "tau_2"])*T_2 + training_load[[i]]
    perf_out[[i]] <- p_0 + params_mat[i, "k_1"]*T_1 - params_mat[i, "k_2"]*T_2
  }
  return(perf_out)
}

```

```

perf_plot <- function(p_0,
                     k_1,
                     tau_1,
                     k_2,
                     tau_2,
                     change_days = NULL,
                     days,
                     training_stim) {
  training_load <- c()
  limit <- 0
  if (training_stim[[1]] == "constant") {
    training_load <- c(rep(training_stim[[2]], days))
    k <- length(k_1)

    # See justification for this below
    limit <- p_0 + training_stim[[2]]*k_1[[k]]/(1-exp(-1/tau_1[[k]])) -
      training_stim[[2]]*k_2[[k]]/(1-exp(-1/tau_2[[k]]))
  }
  tmp_matrix <- params_mat(k_1,
                          tau_1,
                          k_2,
                          tau_2,
                          change_days,
                          days)
  modeled_performance <- perf_tv(p_0, tmp_matrix, training_load)
  tmp_data <- tibble(
    "day" = c(0:days),
    "performance" = c(p_0, modeled_performance),
    "limit" = c(rep(limit, days + 1))
  )

  plot <- ggplot(tmp_data, aes(x = day)) +
    geom_point(aes(y = performance, color = "perf")) +
    geom_line(aes(y = limit, color = "lim"))

  # scale_color_manual("Legend",
  #                     values = c("lim" = "#e31a1c", # this color comes from the theme "Paired"
  #                               "perf" = "black"))
  plot

```

}

Computing the limit of the model

I would like to compute the limit of the predicted performance for the time-invariant model Under the assumption of constant training load. We have

$$p(t) = p_0 + k_1 \sum_{i=1}^{t-1} e^{\frac{t-i}{\tau_1}} w(i) + k_2 \sum_{i=1}^{t-1} e^{\frac{t-i}{\tau_2}} w(i)$$

Assume that $w(i) = C$ for all i . Note that

$$\sum_{i=1}^{t-1} e^{\frac{t-i}{\tau_1}} = e^{1/\tau_1} \sum_{i=1}^{t-1} e^{\frac{(t-1)-i}{\tau_1}} = e^{1/\tau_1} \left(-1 + \sum_{i=0}^{s} (e^{-1/\tau_1})^i \right)$$

Finding the long-run limit of $p(t)$ then amounts to computing

$$\sum_{i=0}^{\infty} (e^{-1/\tau_1})^i = \frac{1}{1 - e^{-1/\tau_1}}$$

Notice that this is a convergent geometric series, $e^{-1/\tau_1} < 1$ when $\tau_1 > 1$ (which we have assumed). Therefore,

$$\sum_{i=1}^{t-1} e^{\frac{t-i}{\tau_1}} = e^{1/\tau_1} \left(-1 + \frac{1}{1 - e^{-1/\tau_1}} \right) = e^{1/\tau_1} \left(\frac{e^{-1/\tau_1}}{1 - e^{-1/\tau_1}} \right) = \frac{1}{1 - e^{-1/\tau_1}}$$

Therefore

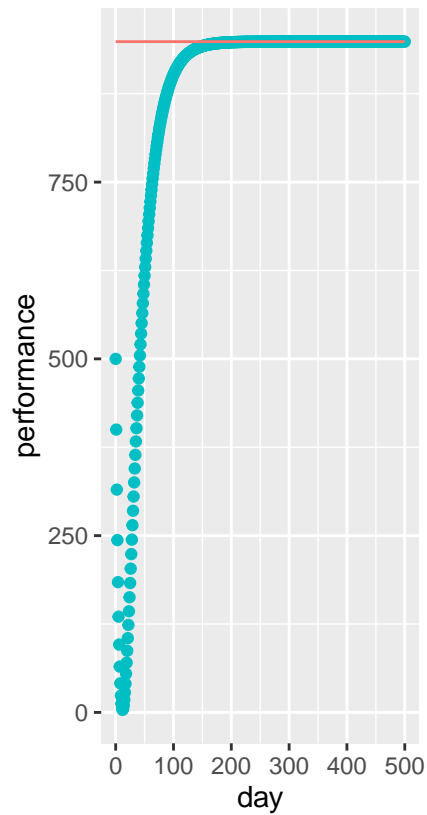
$$\lim_{t \rightarrow \infty} p(t) = p_0 + C \left(\frac{k_1}{1 - e^{-1/\tau_1}} - \frac{k_2}{1 - e^{-1/\tau_2}} \right)$$

This is how we compute the red line in `perf_plot`.

Exploring the time, varying model

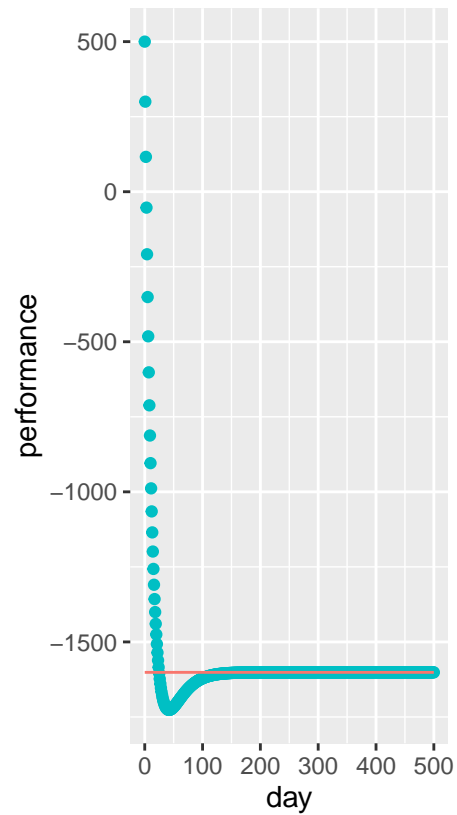
This is a time-invariant plot with our new function, to check that things are working correctly.

```
perf_plot(p_0 = 500,
          k_1 = 1,
          tau_1 = 25,
          k_2 = 2,
          tau_2 = 10,
          days = 500,
          training_stim = list("constant", 100)) +
perf_plot(p_0 = 500,
          k_1 = 2,
          tau_1 = 20,
          k_2 = 4,
          tau_2 = 15,
          days = 500,
          training_stim = list("constant", 100))
```



colour

- lim
- perf

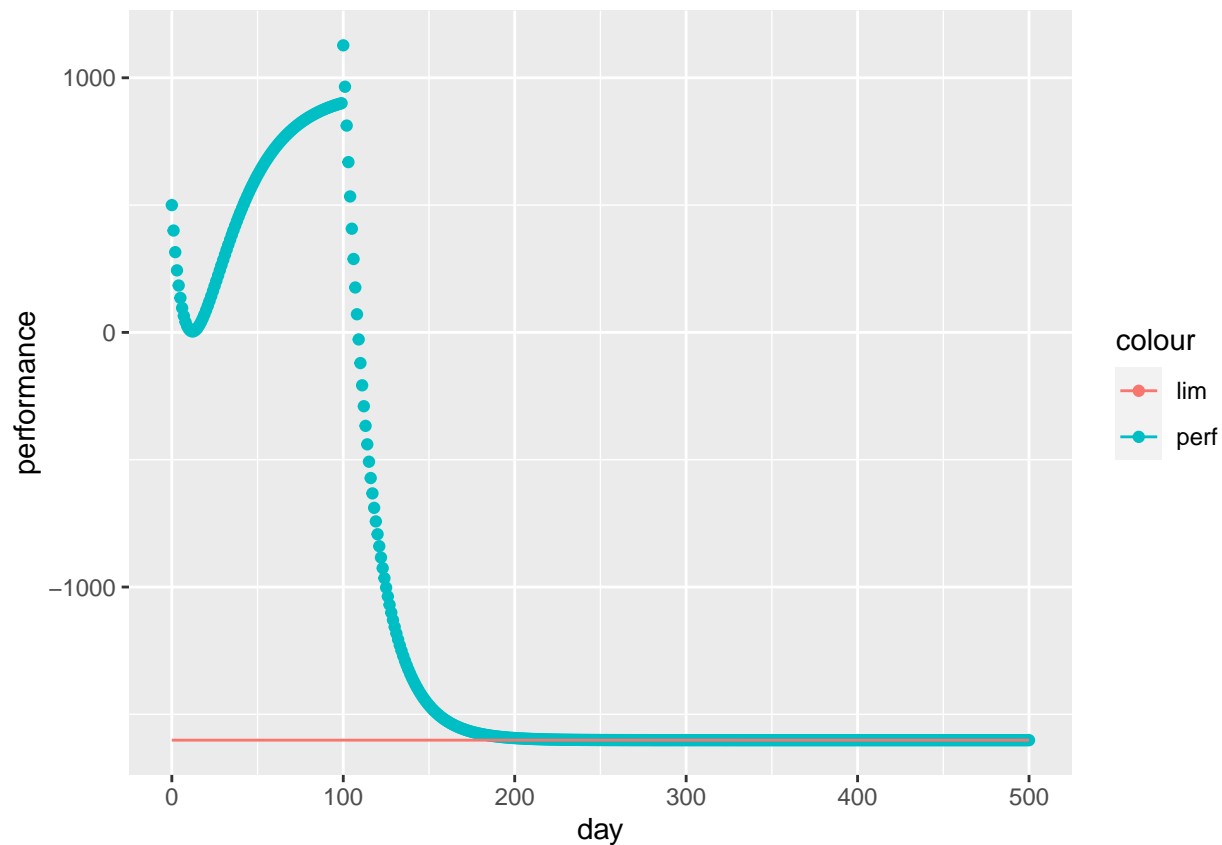


colour

- lim
- perf

Adding one change date

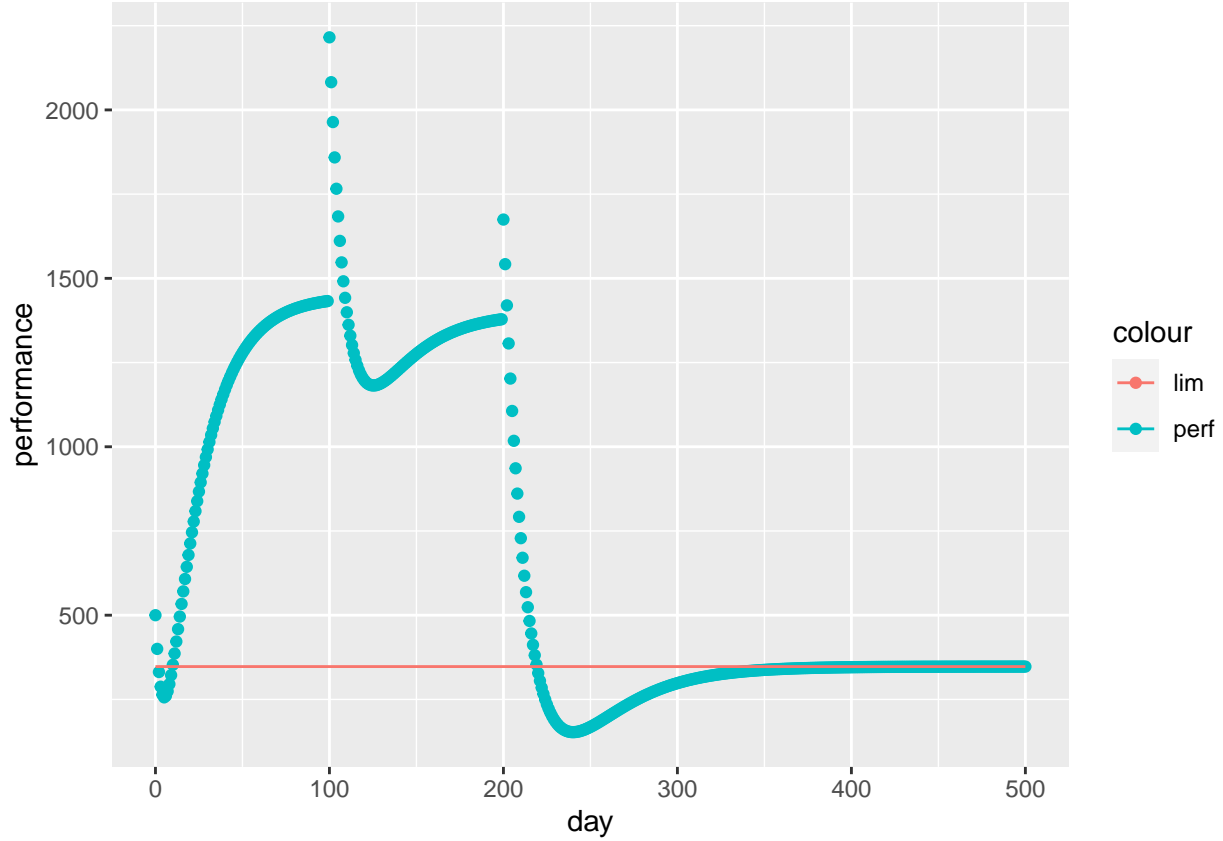
```
perf_plot(p_0 = 500,
          k_1 = c(1, 2),
          tau_1 = c(25, 20),
          k_2 = c(2, 4),
          tau_2 = c(10, 15) ,
          change_days = c(100),
          days = 500,
          training_stim = list("constant", 100))
```



Notice that this doesn't look like the first set of parameters for the first 100 days, and then the other distribution for the rest of the days. Even though it looks like a new curve at the change day, it is reflecting a change of the parameters.

Adding an additional change dage

```
perf_plot(p_0 = 500,
          k_1 = c(1,2,3),
          tau_1 = c(20, 25, 30),
          k_2 = c(2, 4, 6),
          tau_2 = c(5,10, 15),
          change_days = c(100, 200),
          days = 500,
          training_stim = list("constant", 100))
```



To have the effects of an initial negative effect, and a long-run positive effect to a constant stimulus, it seems to be necessary that $k_1 < k_2$ and τ_1 has to be much bigger than

Lets try to use $\frac{k_1}{1-e^{-1/\tau_1}} - \frac{k_2}{1-e^{-1/\tau_2}}$ in the ‘metric’ to evaluate the distance between parameter sets

We can write

$$\text{dist}_{\text{perf}}(z, z') = |f(z) - f(z')|$$

where $f : (\mathbb{R}^2 \times \mathbb{R}_{>0})^2 \rightarrow \mathbb{R}$ is given by

$$z = (k_1, \tau_1, k_2, \tau_2) \mapsto \frac{k_1}{1-e^{-1/\tau_1}} - \frac{k_2}{1-e^{-1/\tau_2}}$$

$\text{dist}_{\text{perf}}$ is a metric, from the properties of the absolute value, except for the fact that the distance between two points has to be positive, since f is not injective.

Since

$$e^{n/x} \approx 1 - \frac{n}{x}$$

we have that

$$1 - e^{-1/x} \approx 1/x$$

Therefore,

$$f(z) = \frac{k_1}{1-e^{-1/\tau_1}} - \frac{k_2}{1-e^{-1/\tau_2}} \approx k_1\tau_1 - k_2\tau_2$$

I don’t believe that this approximation is very good, but it provides a good intuition about this function.

```

lim_func <- function(k_1, tau_1, k_2, tau_2) {
  k_1/(1-exp(-1/tau_1))-k_2/(1-exp(-1/tau_2))
}

params_dist <- function(params_1, params_2) {
  abs(lim_func(params_1[[1]], params_1[[2]], params_1[[3]], params_1[[4]])-
    lim_func(params_2[[1]], params_2[[2]], params_2[[3]], params_2[[4]])
  )
}

```

Applying the “norm” to figure out “distance” from a point

Don’t run this chunk locally

```

params_grid <- expand.grid(k_1 = seq(1,50, length.out=100),
  tau_1 = seq(1,50, length.out=100),
  k_2 = seq(1,50, length.out=100),
  tau_2 = seq(1,50, length.out=100)
)

params_matrix <- as.matrix(params_grid)
dist_vec <- c(rep(0, 100000000))
iter_fn <- function(i, params_matrix) {
  params_dist(params_matrix[i, ], c(1,25,2,10))
}

# parallel computing is a factor for this computation, took a few minutes
dist_vec <- parallel::mcmapply(
  iter_fn,
  i = c(1:100000000),
  MoreArgs = list(params_matrix = params_matrix),
  mc.cores = floor(.9 * parallel::detectCores())
)
# a 100,000,000 element, 100 MB vector
save(dist_vec, file = stringr::str_c(rprojroot::find_rstudio_root_file(),
  "/generated_data/dist_vec.RData"))
# so we don't have to do the computation again

dist_tib <- tibble::tibble(
  "dist" = dist_vec
)

plot_hist <- ggplot(dist_tib, aes(x=dist)) +
  geom_histogram() +
  labs(title = "distance from c(1,25,2,10)" )
plot_hist

# so we can call the outputed plot in this markdownfile, to save time.
ggsave(filename = "dist from single set.pdf",

```

```

plot_hist,
path = stringr::str_c(rprojroot::find_rstudio_root_file(), "/plots"),
device = "pdf")

# this took awhile
cdf_vec <- c(rep(0, 2601))
for (i in 1:2601) {
  cdf[[i]] <- length(which(dist_vec<=i))
}

cdf_data <- tibble(
  "x" = c(1:2601),
  "percent_leq_x" = cdf/100000000
)

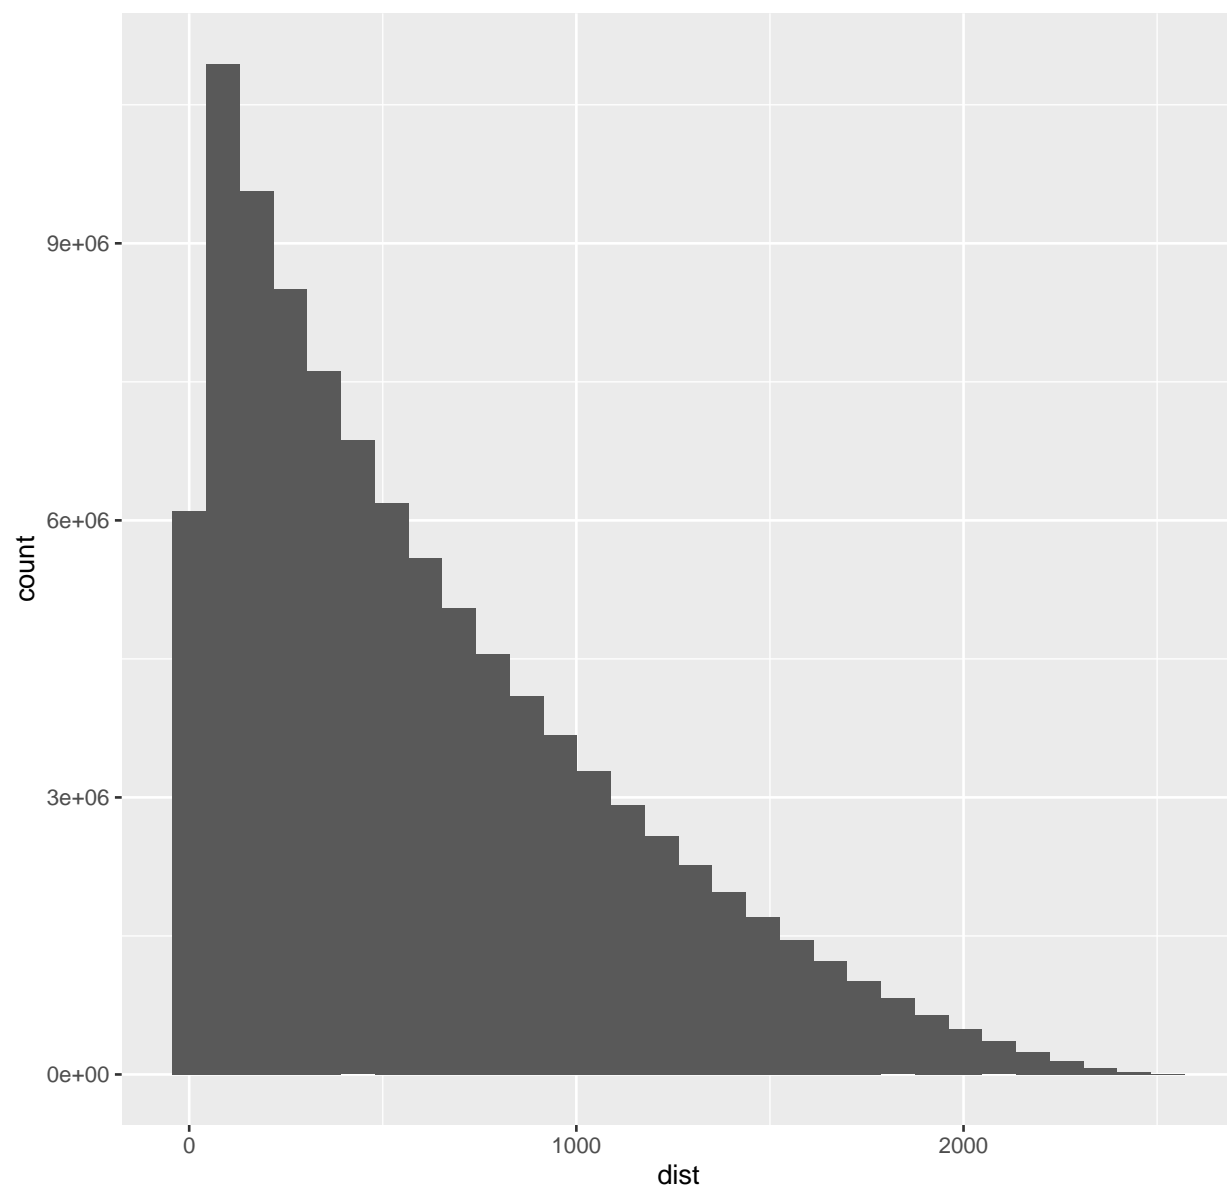
# to save computation time later
save(cdf_vec, file = stringr::str_c(rprojroot::find_rstudio_root_file(),
  "/generated_data/cdf_vec.RData"))

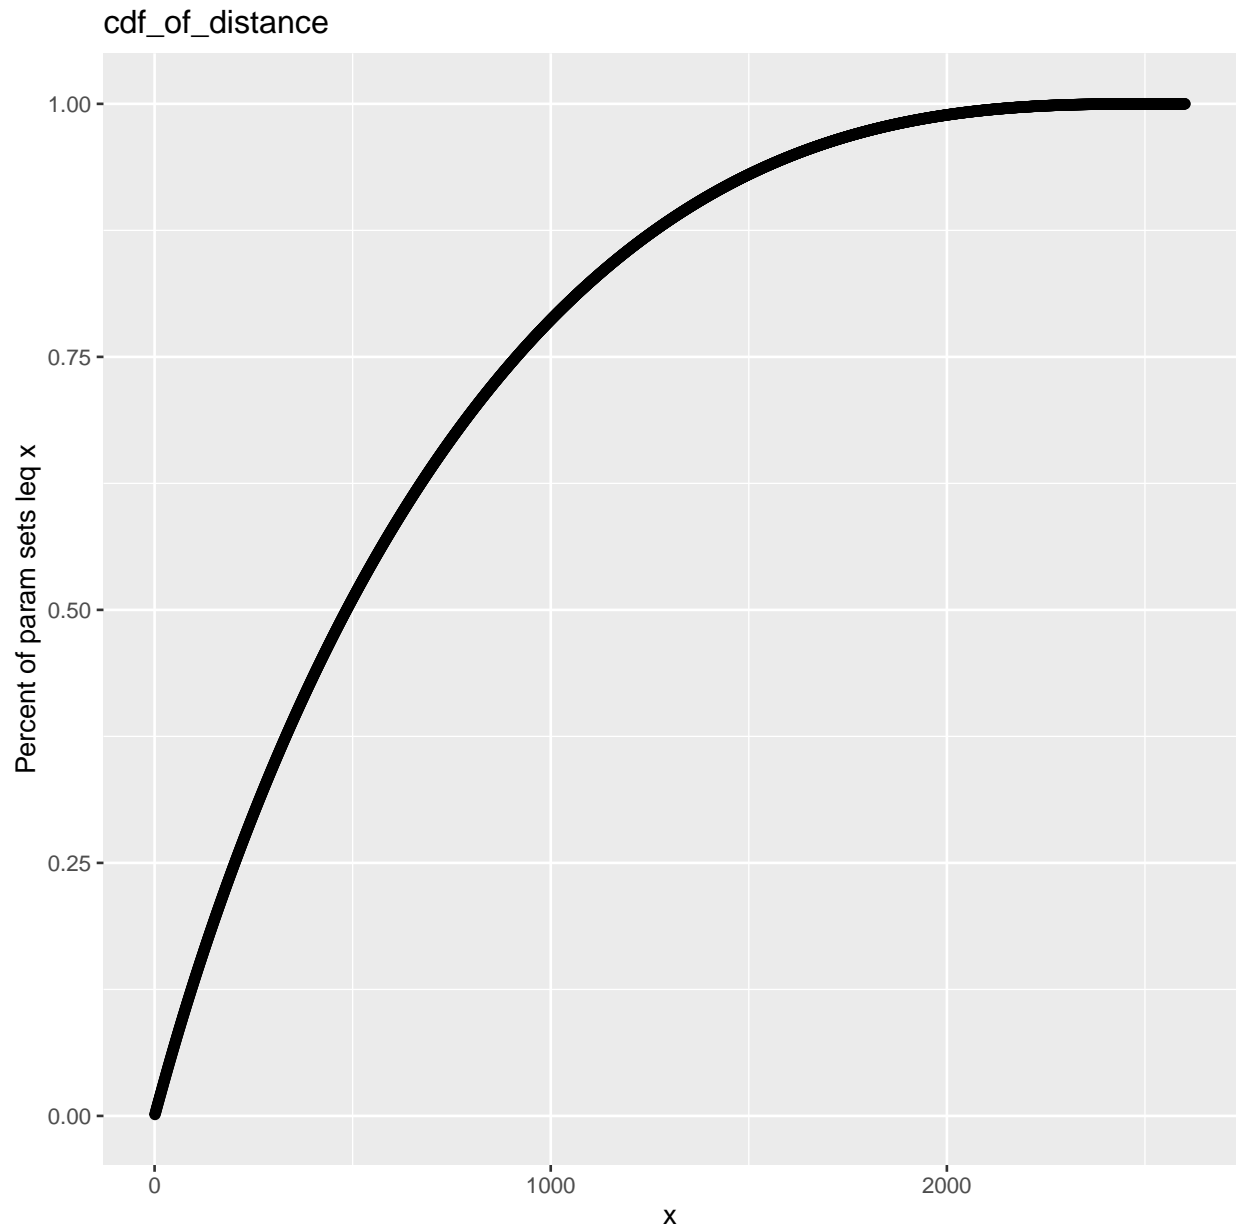
cdf_plot <- ggplot(cdf_data, aes(x=x, y= percent_leq_x)) +
  geom_point() +
  labs(x = "x",
    y = "Percent of param sets leq x",
    title = "cdf_of_distance ")
cdf_plot

ggsave(filename = "cdf_plot.pdf",
  cdf_plot,
  path = stringr::str_c(rprojroot::find_rstudio_root_file(), "/plots"),
  device = "pdf")

```


distance from c(1,25,2,10)





Generating many performance curves with the same limit and plotting

```
set.seed(443)
C_1 <- round(lim_func(1,25,2,10), digits = 1)
n <- 100
vec_1<- runif(n, 5, 50)+C_1
k_1_same <- mapply(function(i) {runif(1, 1, vec_1[[i]])}), c(1:n))
tau_1_same <- -1/log(1-k_1_same/vec_1)

vec_2 <- vec_1-C_1
```

```

k_2_same <- mapply(function(i) {runif(1, 1, vec_2[[i]])}, c(1:n))
tau_2_same <- -1/log(1-k_2_same/vec_2)

day <- 150
p_0 <- 500
training_load <- c(rep(100, day))
same_test_tib <- tibble(
  "day" = c(0:day)
)

for (i in 1:n) {
  same_test_tib[, stringr::str_c("p", i)] <- c(p_0,
  perf_tv(
    p_0 = p_0,
    params_mat = params_mat(
      k_1 = k_1_same[[i]],
      tau_1 = tau_1_same[[i]],
      k_2 = k_2_same[[i]],
      tau_2 = tau_2_same[[i]],
      days = day
    ),
    training_load = training_load
  ))
}

pal_color <- scales::hue_pal()(n)
names(pal_color) <- names(same_test_tib[2:n+1])

cols <- as.list(names(same_test_tib[2:n+1]))
min_or_max_vec <- c(rep(0,n))
for (i in 1:n) {
  if (round(max(same_test_tib[, i+1]), digits = 1)<=950){
    min_or_max_vec[[i]] <- min(same_test_tib[, i+1])
  }
  else {
    min_or_max_vec[[i]] <- max(same_test_tib[, i+1])
  }
}

# reordering the columns so the plot looks nice
ord_min_or_max_vec <- sort(min_or_max_vec)
ord_min_or_max_vec

```

```

## [1] -3602.189461 -3327.611469 -2832.581326 -2432.440204 -2408.966002
## [6] -1648.104660 -975.945776 -918.660903 -626.968981 -572.658665
## [11] -565.992065 -565.048350 -392.771165 -389.042347 -356.617522
## [16] -335.648616 -188.585342 -170.469789 -113.114856 -81.874936
## [21] -28.535047 -3.657609 277.377738 309.045056 332.437935
## [26] 332.500656 340.168229 365.358637 398.443765 403.399607
## [31] 434.195087 446.938957 471.477564 480.863142 500.000000
## [36] 500.000000 500.000000 500.000000 500.000000 500.000000

```

```
## [41] 500.000000 500.000000 500.000000 950.139202 950.904572
## [46] 954.043203 979.344754 1010.548716 1011.162789 1065.364592
## [51] 1091.073545 1104.515909 1126.229670 1129.689630 1160.708664
## [56] 1174.448884 1175.673531 1257.279924 1257.542711 1265.364346
## [61] 1301.476567 1326.798975 1339.303557 1368.451567 1371.699938
## [66] 1454.878860 1501.289621 1528.227053 1542.571439 1586.709490
## [71] 1587.609994 1593.589464 1609.753998 1652.418754 1661.571624
## [76] 1677.399624 1682.020849 1684.326918 1759.961127 1781.524675
## [81] 1848.557754 1934.673381 2031.722969 2049.825880 2076.976777
## [86] 2136.624471 2146.203990 2328.957673 2358.956663 2362.264050
## [91] 2371.730776 2396.775787 2408.549590 2607.199206 3102.314686
## [96] 3617.734852 3675.639829 4100.450575 4302.409166 4916.749802
```

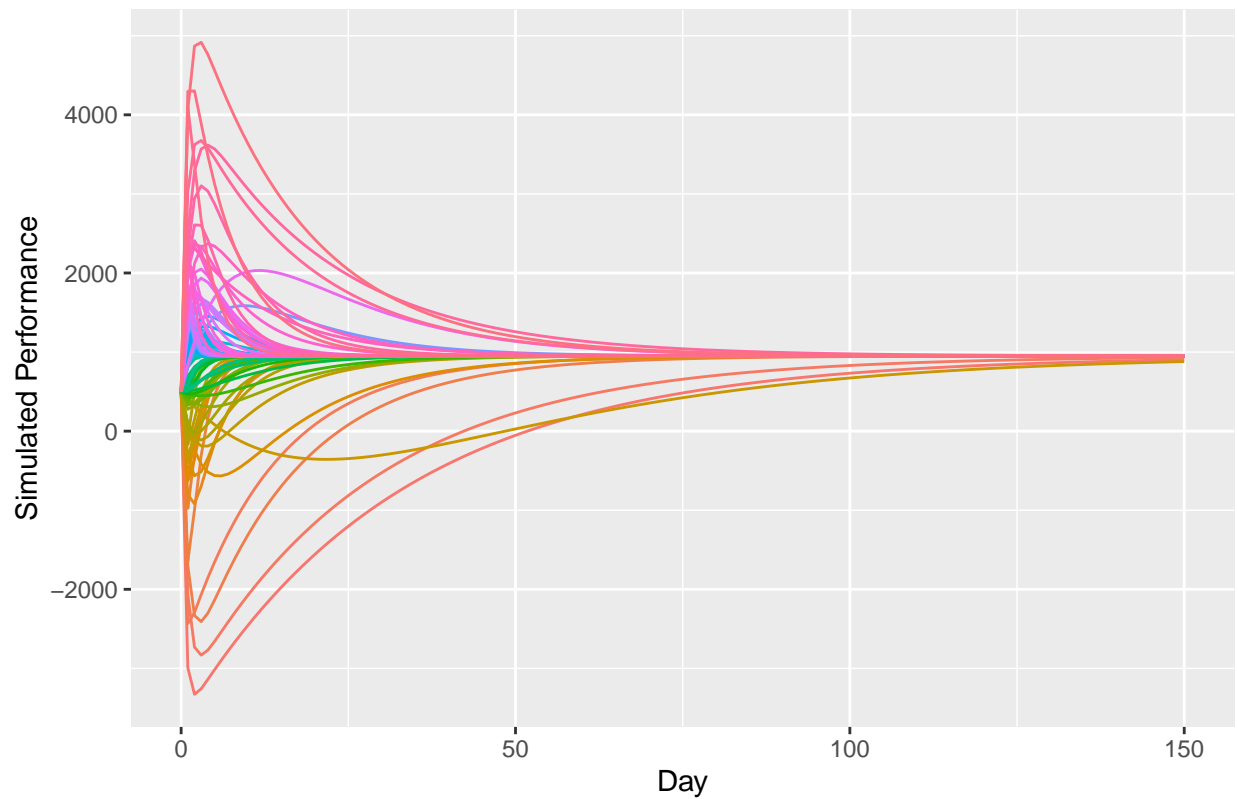
```
permutation_index <- c(rep(NA, n))
for (i in 1:n) {
  value <- min_or_max_vec[[i]]
  output_index <- which(ord_min_or_max_vec == value)
  while(is.na(permutation_index[[i]])) {
    if (output_index[[1]] %in% permutation_index){
      output_index <- output_index[-1]
    }
    else {
      permutation_index[[i]] <- output_index[[1]]
    }
  }
}

same_test_tib_new <- same_test_tib
for (i in 1:n) {
  same_test_tib_new[, permutation_index[[i]] + 1] <- same_test_tib[, i + 1]
}

plot_same <- ggplot(data = same_test_tib_new, aes(x = day)) +
  lapply(cols, function(x) {
    geom_line(aes(y = .data[[x]], color = x))
  }) +
  scale_color_manual(values = pal_color) +
  labs(x = "Day",
       y = "Simulated Performance",
       title = "Simulated Performance of Parameter Sets with Same Limit") +
  theme(legend.position = "none")

plot_same
```

Simulated Performance of Parameter Sets with Same Limit



The colors were sorted by the maximum value. It seems that the maximum value does not determine a convergence rate; there are some curves that peak late, but this generally doesn't happen.

Even though our metric is not truly a metric, it will probably become one when also restrict our curves with minimizing the sum of squared error.